


Fine-tuned encoders beat surface baselines; XLM-R base leads



Encoders: mean of 3 seeds, bars are 95% pair-clustered bootstrap intervals. Baselines: single fit.